

Approximate Leave-One-Out with kernel SVM

Wenda Zhou, Peng Xu

March 6, 2019

1 C-Support Vector Classification

Let \mathbf{K} denote a positive-definite kernel matrix (hence invertible), with $\mathbf{K}_{i,j} = K(\mathbf{x}_i, \mathbf{x}_j)$. The objective of a kernel SVC can be expressed in the “loss + penalty” variational form:

$$\min_{\rho, \alpha} \sum_{j=1}^n \max [0, 1 - y_j h(x_j)] + \frac{\lambda}{2} \alpha^\top \mathbf{K} \alpha, \quad h(x_j) = \mathbf{K}_{:,j}^\top \alpha + \rho. \quad (1)$$

For simplicity we ignore the offset ρ for now. Let S and V be the smooth set and the set of singularities, respectively. For the j -th observation, $j \in S$, we have

$$\dot{\ell}(\mathbf{K}_{:,j}^\top \alpha) = -y_j \cdot \mathbf{1}\{y_j \mathbf{K}_{:,j}^\top \alpha < 1\}, \quad \ddot{\ell}(\mathbf{K}_{:,j}^\top \alpha) = 0.$$

Additionally,

$$\nabla R(\alpha) = \lambda \mathbf{K} \alpha, \quad \nabla^2 R(\alpha) = \lambda \mathbf{K}.$$

Substitute corresponding terms in Thm. 4.1, we deduce the ALO formula for kernel SVC:

$$\mathbf{K}_{:,i}^\top \tilde{\alpha}^i = \mathbf{K}_{:,i}^\top \hat{\alpha} + a_i g_{\ell,i},$$

where

$$a_i = \begin{cases} \frac{1}{\lambda} \mathbf{K}_{:,i}^\top \left[\mathbf{K}^{-1} - \mathbf{K}^{-1} \mathbf{K}_{:,V} \left(\mathbf{K}_{:,V}^\top \mathbf{K}^{-1} \mathbf{K}_{:,V} \right)^{-1} \mathbf{K}_{:,V}^\top \mathbf{K}^{-1} \right] \mathbf{K}_{:,i} & i \in S, \\ \left[\lambda \left(\mathbf{K}_{:,V}^\top \mathbf{K}^{-1} \mathbf{K}_{:,V} \right)^{-1} \right]^{-1}_{ii} & i \in V, \end{cases}$$

and

$$g_{\ell,S} = -y_S \odot \mathbf{1}\{y_S \mathbf{K}_{:,S}^\top \alpha < 1\}, \quad g_{\ell,V} = \left(\mathbf{K}_{:,V}^\top \mathbf{K}_{:,V} \right)^{-1} \mathbf{K}_{:,V}^\top \left[\sum_{j \in S: y_j \mathbf{K}_{:,j}^\top \alpha < 1} y_j \mathbf{K}_{:,j} - \lambda \mathbf{K} \alpha \right].$$

2 ε -Support Vector Regression

The ε -SVR is associated with the ε -insensitive loss function. Its objective function can be written as:

$$\min_{\rho, \alpha} \sum_{j=1}^n \max [0, |y_j - h(x_j)| - \varepsilon] + \frac{\lambda}{2} \alpha^\top K \alpha, \quad h(x_j) = K_{:,j}^\top \alpha + \rho. \quad (2)$$

So for the j -th observation, $j \in S$, we have

$$\dot{\ell}(K_{:,j}^\top \alpha) = -\text{sgn} \left(K_{:,j}^\top \alpha \right) \cdot \mathbf{1} \left\{ \left| y_j - K_{:,j}^\top \alpha \right| \geq \varepsilon \right\}, \quad \ddot{\ell}(K_{:,j}^\top \alpha) = 0.$$

Our ALO recipe will then be exactly the same as in C-SVC, except

$$g_{\ell,S} = -\text{sgn} \left(K_{:,S}^\top \alpha \right) \odot \mathbf{1} \left\{ \left| y_j - K_{:,S}^\top \alpha \right| \geq \varepsilon \right\},$$

and

$$g_{\ell,V} = \left(K_{:,V}^\top K_{:,V} \right)^{-1} K_{:,V}^\top \left[\sum_{j \in S: |y_j - K_{:,j}^\top \alpha| \geq \varepsilon} \text{sgn} \left(K_{:,j}^\top \alpha \right) K_{:,j} - \lambda K \alpha \right].$$

3 ν -SVC and ν -SVR

ν -SVC and ν -SVR are introduced in [Sch+00], with $\nu \in (0, 1]$ a new parameter that can control the number of support vectors. Notably, ν replaces C in C-SVC and ε in ε -SVR, and in the latter case ε itself becomes an optimization argument.

Because of the following equivalencies, there is no need to figure out the variational problem in order to derive ALO for ν -SVC and ν -SVR:

- Let α^* be the dual solution of ν -SVC with parameters (C, ν) , then α^*/ρ^* is an optimal solution of C-SVC with $C = 1/(n\rho^*)$. See [CL01].
- Let α^* be the dual solution of ν -SVR with parameters (C, ν) , then α^* is also the solution of ε -SVR with parameters $(C/n, \varepsilon^*)$. See [CL02].

Here ρ^* and ε^* are the corresponding primal solutions, and can be worked out quickly from α^* through equations derived from KKT conditions. See [CL11, Section 4.2].

4 Multiclass-classifier

To handle the multinomial classification, LIBSVM adopts the one-versus-one strategy: for a dataset with $y \in \mathcal{S}$, $|\mathcal{S}| = K$, LIBSVM will train $K(K-1)/2$ classifiers $\hat{f}_{(u,v)}(\mathbf{x})$, one for each 2-combination (u, v) of \mathcal{S} . The prediction for a new observation \mathbf{x}_{new} is then obtained through voting: $\hat{y}_{\text{new}} = \text{mode}\{\hat{f}_{(u,v)}(\mathbf{x}_{\text{new}})\}$, and the label with the lowest index will be chosen in case of ties.

Note that for a specific pair (u', v') , the leave- i -out prediction $\hat{f}_{(u',v')}^{\setminus i}(\mathbf{x}_j)$ is exactly the full-data prediction $\hat{f}_{(u',v')}(\mathbf{x}_j)$ when y_i is neither u' nor v' , for \mathbf{x}_i does not contribute to the training of

$f_{(u',v')}(\cdot)$ in anyway. Therefore, to obtain the multiclass ALO prediction $\tilde{y}_j^{\setminus i}$, we first construct all $K(K-1)/2$ full-data predictions, then replace the $K-1$ ones whose training are affected by y_i with the corresponding binary ALO predictions, and finally ensemble the result through popular voting.

5 Computation

To compute ALO for SVM, we must perform a variety of operations with the kernel matrix $\mathbf{K} \in \mathbb{R}^{n \times n}$. In practice n can often be quite large, therefore, naively implement the ALO formula will result in severe loss of efficiency. This issue can be, however, alleviated by exploiting the symmetry of \mathbf{K} : consider the Cholesky decomposition $\mathbf{K} = \mathbf{L}\mathbf{L}^\top$, then for an arbitrary index set $E \subseteq \{1, \dots, n\}$, we have

$$\mathbf{L}^{-1}\mathbf{K}_{:,E} = \mathbf{L}^{-1}\mathbf{L}\mathbf{L}_{E,:}^\top = \mathbf{L}_{E,:}^\top.$$

By further considering the thin QR factorization $\mathbf{L}_{V,:}^\top = \mathbf{Q}_V\mathbf{R}_V$, where \mathbf{Q}_V semi-orthogonal and \mathbf{R}_V upper-triangular, we may expressed a_s , $s \in S$, as a difference of two quadratic form:

$$\begin{aligned} \lambda a_s &= \mathbf{K}_{:,s}^\top \mathbf{K}^{-1} \mathbf{K}_{:,s} - \mathbf{K}_{:,s}^\top \mathbf{K}^{-1} \mathbf{K}_{:,V} \left(\mathbf{K}_{:,V}^\top \mathbf{K}^{-1} \mathbf{K}_{:,V} \right)^{-1} \mathbf{K}_{:,V}^\top \mathbf{K}^{-1} \mathbf{K}_{:,s} \\ &= \mathbf{K}_{:,s}^\top (\mathbf{L}\mathbf{L}^\top)^{-1} \mathbf{K}_{:,s} - \mathbf{K}_{:,s}^\top (\mathbf{L}\mathbf{L}^\top)^{-1} \mathbf{K}_{:,V} \left(\mathbf{K}_{:,V}^\top (\mathbf{L}\mathbf{L}^\top)^{-1} \mathbf{K}_{:,V} \right)^{-1} \mathbf{K}_{:,V}^\top (\mathbf{L}\mathbf{L}^\top)^{-1} \mathbf{K}_{:,s} \\ &= (\mathbf{L}^{-1} \mathbf{K}_{:,s})^\top \mathbf{L}^{-1} \mathbf{K}_{:,s} - (\mathbf{L}^{-1} \mathbf{K}_{:,s})^\top (\mathbf{L}^{-1} \mathbf{K}_{:,V}) \left[(\mathbf{L}^{-1} \mathbf{K}_{:,V})^\top (\mathbf{L}^{-1} \mathbf{K}_{:,V}) \right]^{-1} (\mathbf{L}^{-1} \mathbf{K}_{:,V})^\top (\mathbf{L}^{-1} \mathbf{K}_{:,s}) \\ &= \mathbf{L}_{s,:} \mathbf{L}_{s,:}^\top - \mathbf{L}_{s,:} \mathbf{L}_{V,:}^\top \left(\mathbf{L}_{V,:} \mathbf{L}_{V,:}^\top \right)^{-1} \mathbf{L}_{V,:} \mathbf{L}_{s,:}^\top \\ &= \mathbf{L}_{s,:} \mathbf{L}_{s,:}^\top - (\mathbf{Q}_V^\top \mathbf{L}_{s,:})^\top (\mathbf{Q}_V^\top \mathbf{L}_{s,:}). \end{aligned}$$

Similarly, a_v with $v \in V$ can be evaluated through the inner product of two vectors:

$$(\lambda a_v)^{-1} = \left(\mathbf{L}_{V,:} \mathbf{L}_{V,:}^\top \right)_{v,v}^{-1} = (\mathbf{R}_V^\top \mathbf{Q}_V^\top \mathbf{Q}_V \mathbf{R}_V)_{v,v}^{-1} = \mathbf{R}_{v,:}^{-1} \mathbf{R}_{v,:}^{-1}.$$

As a result, with \mathbf{L} , and thus $\mathbf{L}_{S,:}$, $\mathbf{L}_{V,:}$ computed beforehand, we avoided direct multiplication of large matrices and reduced computational costs.

References

- [CL01] Chih-Chung Chang and Chih-Jen Lin. "Training v-Support Vector Classifiers: Theory and Algorithms". In: *Neural Computation* 13.9 (2001), pp. 2119–2147. DOI: 10.1162/089976601750399335.
- [CL02] Chih-Chung Chang and Chih-Jen Lin. "Training v-Support Vector Regression: Theory and Algorithms". In: *Neural Computation* 14.8 (2002), pp. 1959–1977. DOI: 10.1162/089976602760128081.
- [CL11] Chih-Chung Chang and Chih-Jen Lin. "LIBSVM: A Library for Support Vector Machines". In: *ACM Trans. Intell. Syst. Technol.* 2.3 (May 2011), 27:1–27:27. ISSN: 2157-6904. DOI: 10.1145/1961189.1961199. URL: <https://www.csie.ntu.edu.tw/~cjlin/papers/libsvm.pdf>.

- [Sch+00] Bernhard Schölkopf et al. “New Support Vector Algorithms”. In: *Neural Computation* 12.5 (2000), pp. 1207–1245. doi: 10.1162/089976600300015565.